The Problem

BERT is powerful, but too large and expensive for efficient real-world deployment.

1 Large Model

- ~110M parameters
- High memory footprint

2 Slow & Costly

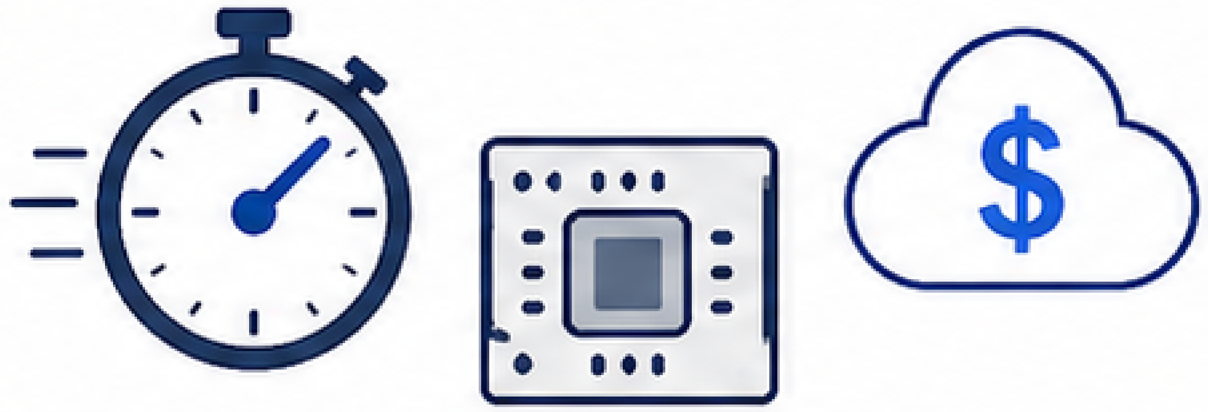

- Slower inference
- More compute required

3 Harder to Deploy

- Less practical for mobile, edge, and low-latency use

Challenge: Keep most of BERT's performance while making the model smaller, faster, and cheaper.

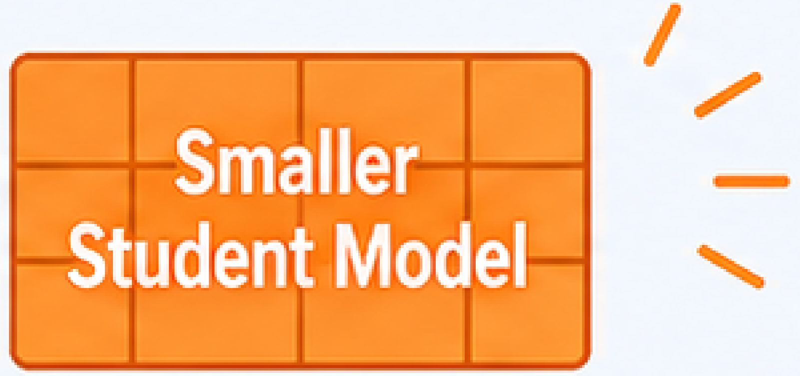

This is the motivation behind DistilBERT: compress BERT without losing much accuracy.